

Task: Linear Regression with Custom Implementation and Benchmarking

You are required to implement and evaluate a simple linear regression model using gradient descent. Follow the steps below to complete the task:

1. Implement a Cost Function

- Create a function to compute the Mean Squared Error (MSE), which will be used to evaluate the performance of the model.

2. Implement Gradient Descent for Linear Regression

- Develop a gradient descent algorithm to optimize the weights of a linear regression model.
- Use the provided dataset for training.
- Data Preprocessing:
 - Normalize the input data using standard scaling or another suitable normalization technique.
- Visualization: Plot the normalized datapoints before training the model.

3. Evaluate Your Custom Model

- After training, measure the performance of your model using the following metrics:
 - o Mean Squared Error (MSE)
 - o Root Mean Squared Error (RMSE)
 - o Mean Absolute Error (MAE)
 - o R-squared (R-squared score)

4. Compare with Scikit-learn Implementation

- Fit a linear regression model using `sklearn.linear_model.LinearRegression` on the same preprocessed dataset.
- Retrieve and report the `coef_` (model coefficients) and `intercept_` (bias term).
- Compare these values with the parameters obtained from your custom implementation. Discuss any similarities or differences.

5. Implement a piecewise regression model on the provided dataset

- Find the suitable number of breakpoints
- Evaluate the piecewise regression model and compare it to the previous implementation