



Sequence modeling

Hossam Ahmed

Ziad Waleed

Mario Mamdouh

Sequence data



Sequence data is data where the **order** of elements is meaningful, and each element depends on previous (and sometimes future) elements in the sequence.



Video



Audio



Time-series



Text



DNA/RNA

IEEE ML S25' training sessions



GPS trajectory

Sequence data features



Order matters : The position of each element is essential; reordering changes the meaning



Context dependence : Each element is related to previous (and sometimes future) elements



Variable length : Sequences can be of different lengths — not all inputs have the same size.

Sequence modeling challenges

Sequence modeling requires handling **variable-length inputs**, preserving **temporal dependencies** and order, **aligning inputs and outputs over time**, and maintaining **long-term memory** — all of which are difficult for traditional neural networks.


Model Type	Limitation
ANNs	<ul style="list-style-type: none">• Fixed input size: can't handle variable-length sequences.• No memory: each input is treated independently.• No temporal awareness.
CNNs	<ul style="list-style-type: none">• Local receptive fields only capture short-range patterns.• Work with grid data and don't expect a sequence.• Position is learned indirectly (via filters).• Can't model long-term dependencies well.

Sequence modeling types

Many to One: Sequence is reduced to a single output — e.g., classifying a full sentence or time-series


One to Many: Fixed input produces a sequence — e.g., image → caption generation

Many-to-many: maps input to output sequences, either **aligned** (same length, e.g., tagging) or **unaligned** (different length, e.g., translation).



Recurrent Neural Networks (RNN)

Recurrent Neural Networks (RNNs) are neural architectures designed for sequential data, where the model maintains a **hidden state** that is updated at each time step to capture **context and dependencies** across the sequence.




Inference using RNNs

At each time step t , the RNN :

1. Takes the **current input** X_t and the **previous hidden state** h_{t-1}
2. Combine them using learned weight matrices (W, U)
3. Applies a non-linear activation function to **update** the hidden state h_t
4. Uses the updated hidden state to produce the output y_t

- The hidden state acts like **memory**, carrying information from previous time steps forward.




$$h_t = g(Uh_{t-1} + WX_t)$$

$$y_t = f(Vh_t)$$


Inference using RNNs

- **Input vector** $X_t \in \mathbb{R}^{n_x}$, is the number of features per time step
 - Then $W \in \mathbb{R}^{n_h \times n_x}$
- **Hidden state** $h_t \in \mathbb{R}^{n_h}$, number of hidden units (memory capacity) like the neurons in ANNs.
 - Then $U \in \mathbb{R}^{n_h \times n_h}$
- **Output vector** $y_t \in \mathbb{R}^{n_y}$, task dependent
 - Then $V \in \mathbb{R}^{n_h \times n_y}$




Inference using RNNs

- h_0 is initialized to be zeros as there is no previous hidden states.
- Weight matrices W , U , V are shared across all time steps
- This is many to many RNN
- <https://joshvarty.github.io/VisualizingRNNs/>



$$h_t = g(Uh_{t-1} + Wx_t)$$

$$y_t = f(Vh_t)$$



Limitations of RNNs

Vanishing/Exploding gradients: Gradients shrink or grow exponentially during backpropagation through time, making training unstable or slow.

Short-term memory: Struggles to capture long-range dependencies — important information fades over time.


Sequential computation: Cannot parallelize across time steps — slows training and inference.

Inventions to alleviate RNNs limitations

Invention	Solves	How It Helps
LSTM (Long Short-Term Memory)	Vanishing gradients, long-term memory	Adds gates (forget, input, output) and a cell state to maintain long-term dependencies
GRU (Gated Recurrent Unit)	Like LSTM, but with fewer parameters	Combines forget & input gates into an update gate — simpler, faster training
Bidirectional RNNs	Limited context, one-directional dependency	Processes sequences forward and backward , improving context awareness
Attention Mechanisms	Long-range dependency, fixed memory bottleneck	Allows the model to focus selectively on relevant parts of the sequence
Transformer Architecture	Sequential computation, long-term memory	Replaces recurrence with self-attention , enabling parallelization and better long-range modeling

Long Short-Term memory (LSTM)


- LSTM was introduced to solve the **vanishing gradients** problem to be able to train deeper RNN and to alleviate the loss of old information in the sequence.



Long Short-Term memory (LSTM)

An **LSTM** unit consists of a **cell state** C_t , a **hidden state** h_t , and **three gates** (forget, input, output) that **control the flow of information** to preserve important signals across long sequences.

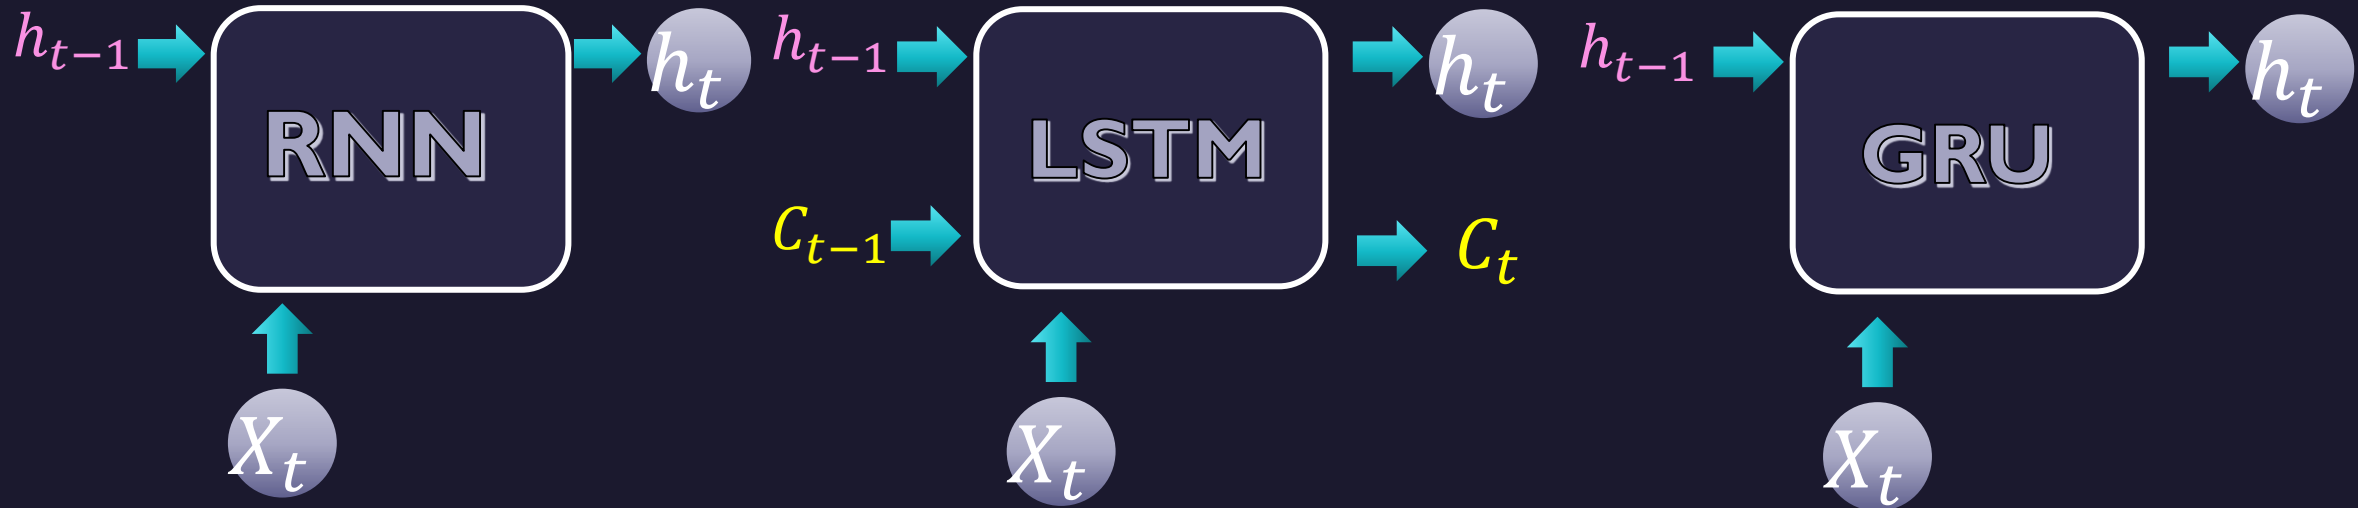
- Forget gate f_t discards irrelevant past memory
- Input gate i_t decides what new info to store
- Cell state C_t combining information from old and new memory to update the long memory
- Output gate o_t use the updated memory to calculate the next hidden state



Long Short-Term memory (LSTM)

Gate	Function	Equation
Forget Gate	Discards irrelevant past memory	$f_t = \sigma(\mathbf{U}_f \mathbf{h}_{t-1} + \mathbf{W}_f \mathbf{X}_t + \mathbf{b}_f)$ $\mathbf{K}_t = \mathbf{C}_{t-1} \odot f_t$
Input Gate	Decides what new info to store	$\mathbf{i}_t = \sigma(\mathbf{U}_i \mathbf{h}_{t-1} + \mathbf{W}_i \mathbf{X}_t + \mathbf{b}_i)$ $\mathbf{g}_t = \tanh(\mathbf{U}_g \mathbf{h}_{t-1} + \mathbf{W}_g \mathbf{X}_t)$ $\mathbf{J}_t = \mathbf{i}_t \odot \mathbf{g}_t$
Cell Update	Combines old and new memory	$\mathbf{C}_t = \mathbf{C}_{t-1} \odot f_t + \mathbf{i}_t \odot \mathbf{g}_t$ $\mathbf{C}_t = \mathbf{K}_t \odot \mathbf{J}_t$
Output Gate	Controls what memory is revealed	$\mathbf{o}_t = \sigma(\mathbf{U}_o \mathbf{h}_{t-1} + \mathbf{W}_o \mathbf{X}_t + \mathbf{b}_o)$ $\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{C}_t)$

RNN vs LSTM vs GRU



RNN has **no gates** and the **fewest parameters**, making it fast but weak at remembering. **LSTM** uses **3 gates** (input, forget, output) and a **cell state**, resulting in **4× the parameters** of a simple RNN, **GRU** uses **2 gates** (update and reset), merges the cell and hidden state, and needs **3× the parameters** of an RNN.

Number of parameters

- To calculate the number of parameters for a recurrent neural network we use this formula
 - $W * |GATES| + U * |GATES| + b * |GATES|$
 - $W \in \mathbb{R}^{n_h \times n_x}, U \in \mathbb{R}^{n_h \times n_h}$
 - Assuming it's Many-to-One (e.g. Classification) so we don't bother with V
- If we have 64 recurrent unit, and the input vector resulting from the embedding layer 45
 - For Simple RNN $(45 * 64) * |1| + (64 * 64) * |1| + 64 * |1| = 7040$
 - For LSTM $(45 * 64) * |4| + (64 * 64) * |4| + 64 * |4| = 28160$
 - For GRU $(45 * 64) * |3| + (64 * 64) * |3| + 64 * |3| = 21120$
 - In TensorFlow GRUs has two biases for each recurrent unit so
 - $(45 * 64) * |3| + (64 * 64) * |3| + 64 * 2 * |3| = 21312$

References

- <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-recurrent-neural-networks>
- https://d2l.ai/chapter_recurrent-modern/bi-rnn.html
- <https://youtu.be/AsNTP8Kwu80?si=dDL6yuahwIzocxIC>
- <https://youtu.be/YCzL96nL7j0?si=ZerUz-cTqG-EwMvb>

See

- <https://joshvarty.github.io/VisualizingRNNs/> ✨
- <https://distill.pub/2019/memorization-in-rnns/>
- <https://damien0x0023.github.io/rnnExplainer/> ✨