



# Transformers


*Hossam Ahmed*

*Ziad Waleed*

*Mario Mamdouh*




In the old days of language modeling, the throne was long held by **Recurrent Neural Networks** and **convolution-based architectures**



Everything changed in **2017** with the **rise of transformers** that leverages **attention mechanism**


# Introduction

- RNN architectures were the most used model for dealing with **sequential data**
- RNNs function similarly to a feed-forward neural network but process the input sequentially, **one element at a time**.



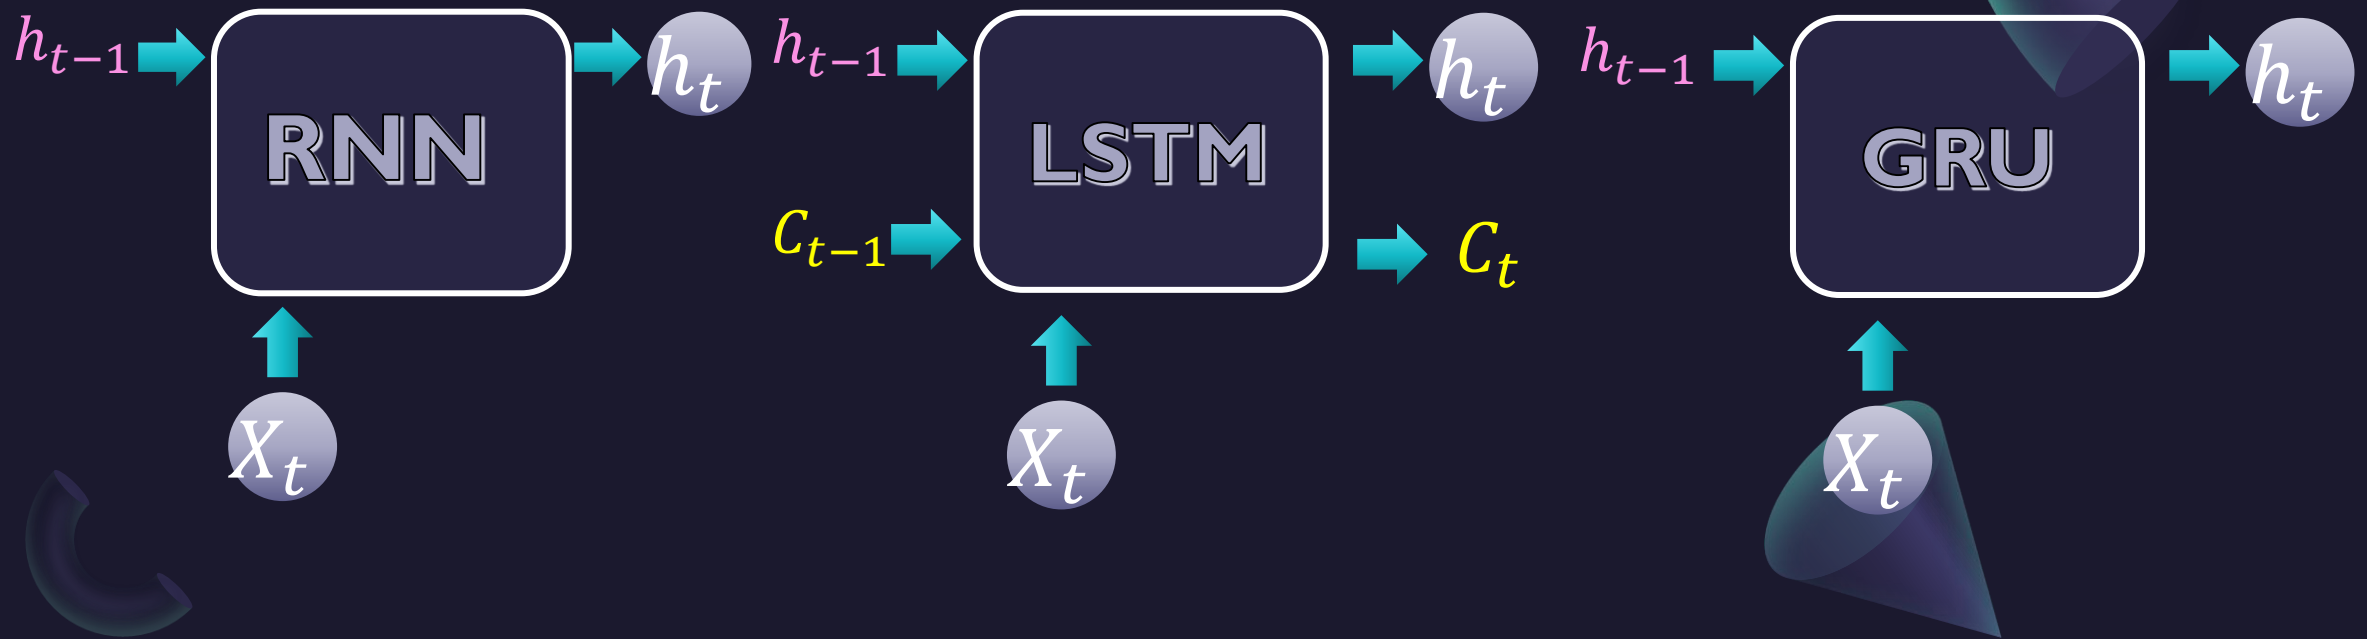
# Introduction

- LSTM was introduced to solve the **vanishing gradients** problem to be able to train more deep RNN and to alleviate the loss of old information in the sequence.



# Introduction

- GRU introduced in 2014, omitting context vector, resulting in a fewer parameters
- **Making us able to create deeper models with fewer parameters and faster training**






# Introduction **problems with RNNs**


- **Hard to parallelize** as they process the data sequentially, one input after the other so doesn't make use of modern GPUs.
- **Difficulty with Long-Term Dependencies** this is due to the vanishing gradients problem that can cause loss of information when the chain of RNN units grows.
- **Limited Context Understanding** RNNs have a fixed-size context window determined by the length of the sequence they process.

# Introduction **Common RNN NLP Architectures**



**Sequence Labeling**


**Named Entity Tagging**



**Sequence Classification**

**Sentiment analysis**

# Introduction **Common RNN NLP Architectures**



**Language modeling**

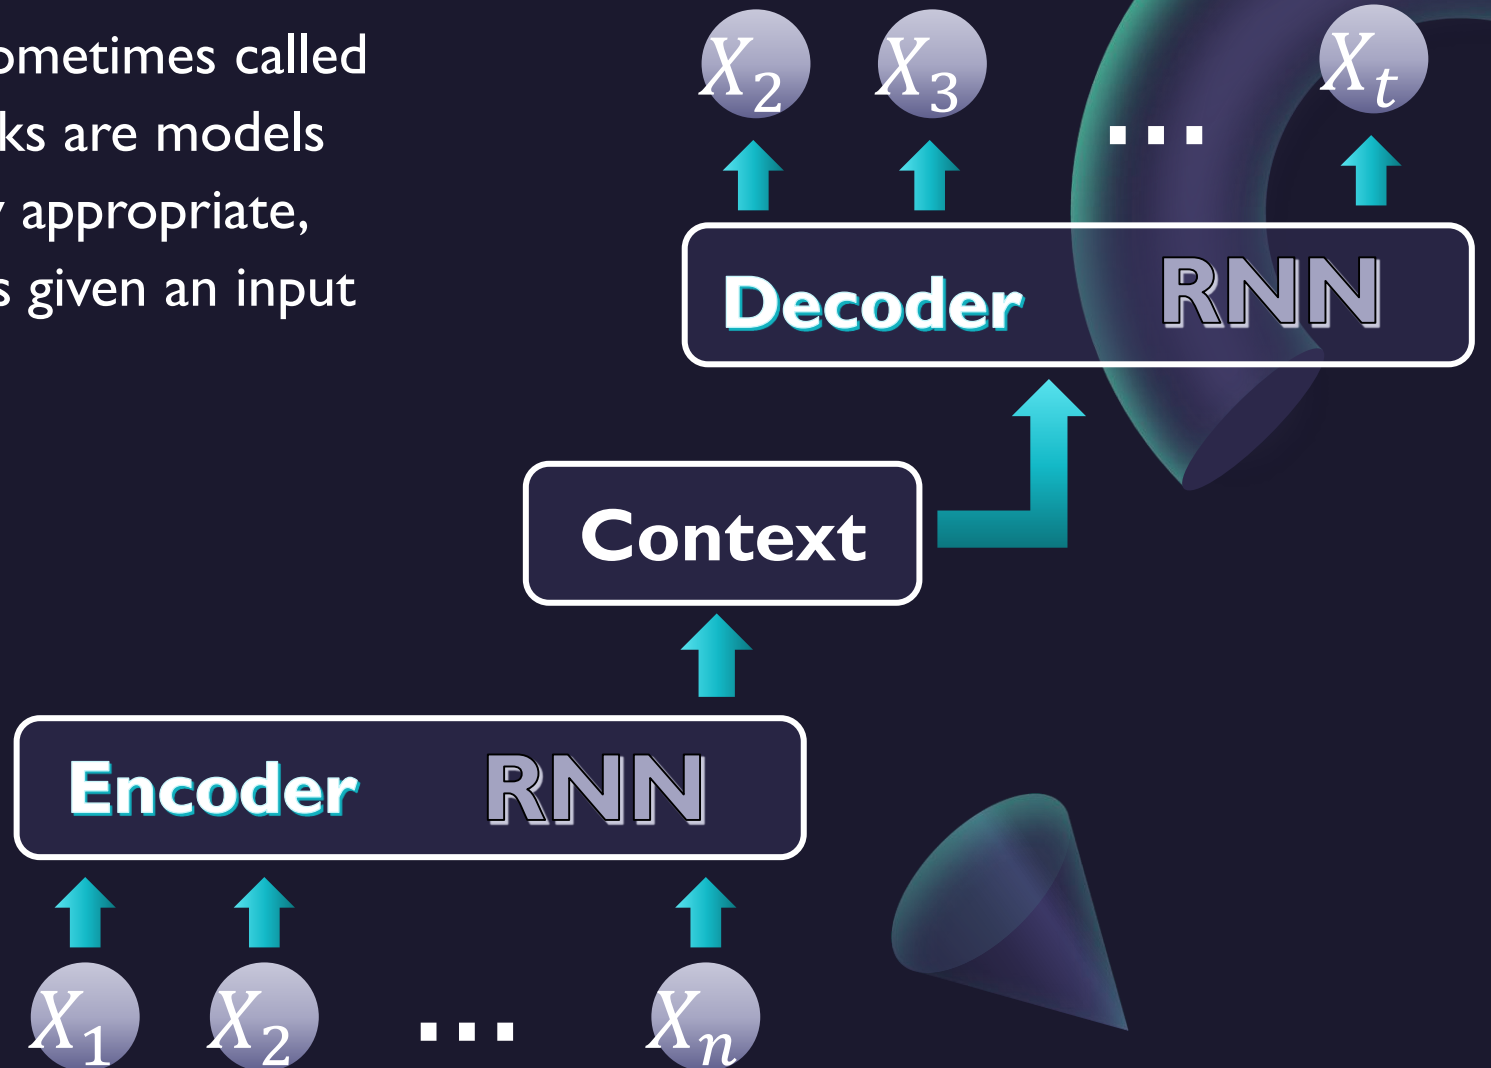
**What is the next word...?**



# Introduction **Common RNN NLP Architectures**


- **Encoder-decoder networks**, sometimes called **sequence-to-sequence** networks are models capable of generating contextually appropriate, arbitrary length, output sequences given an input sequence.

## Encoder-Decoder Translation



# Introduction **context bottleneck**

- Requiring the context to be **the only encoder's final hidden state** forces all the information from the entire source sentence to pass through this representation bottleneck.
- **Bottleneck because**
  - it must represent absolutely everything about the meaning of the source text
  - since the **Decoder knows only the context vector** in this bottleneck

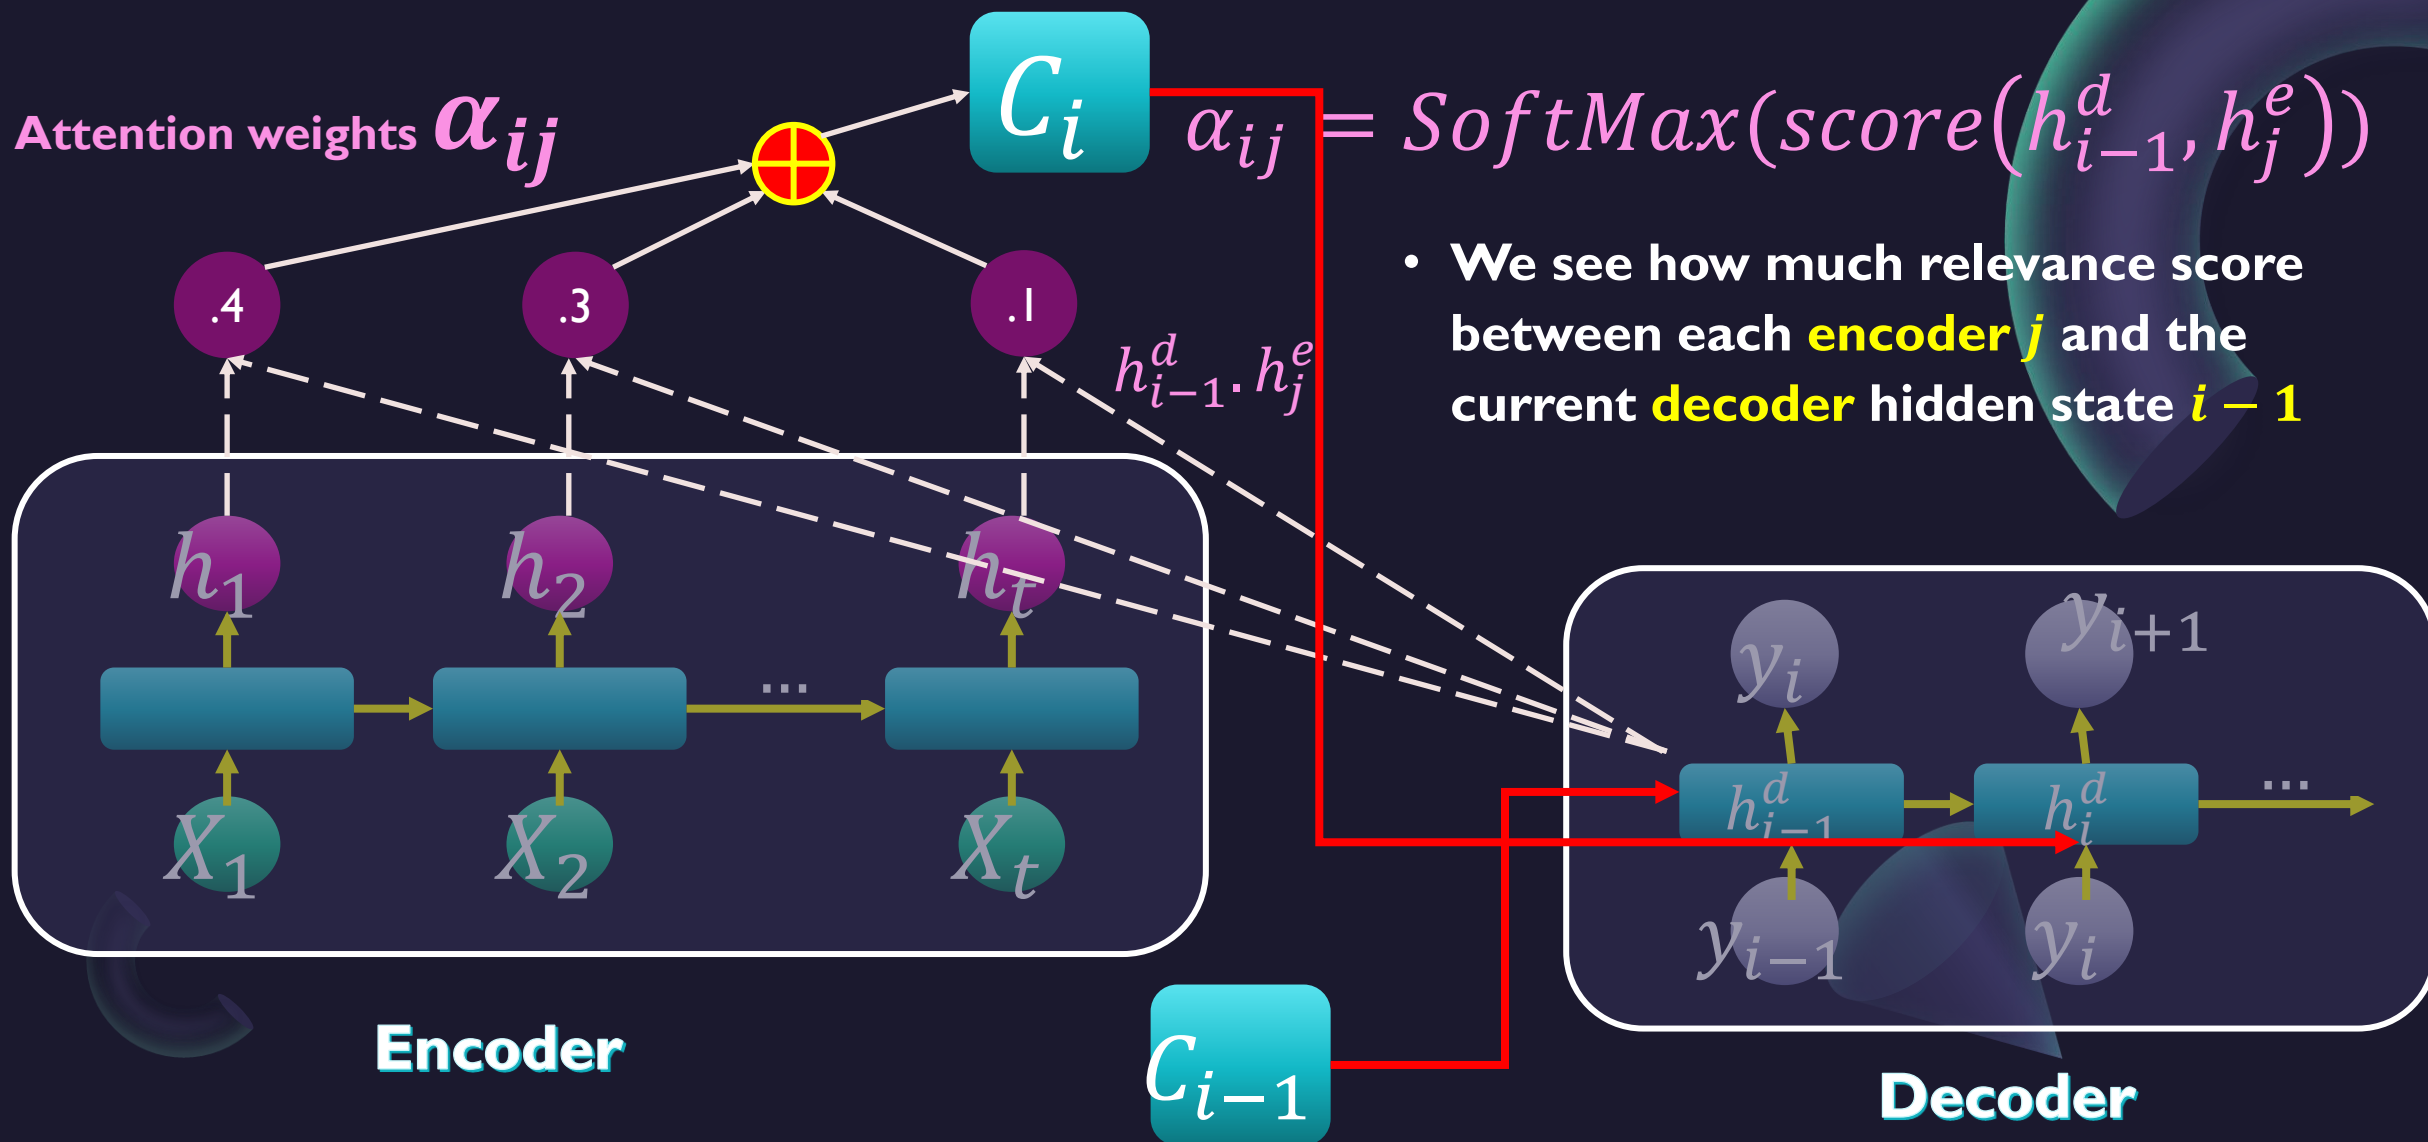


# Attention is all you need **attention**

- **attention mechanism** is a solution to the bottleneck problem, a way of **allowing the decoder to get information from all the hidden states** of the **encoder**, not just the last hidden state.
- The **idea of attention** is instead to **create the single fixed-length vector  $c$**  by taking a weighted sum of all the encoder hidden states.
  - The weights focus on ('attend to') a particular part of the source text that is relevant for the token the decoder is currently producing.
  - Attention thus **replaces the static context vector** with one that is **dynamically derived from the encoder hidden states, different for each token in decoding**

# Attention is all you need **attention**


$$C_i = \sum \alpha_{ij} \cdot h_j^e$$






# Attention is all you need **Self-Attention**

- **Self-Attention** can be thought of a way to **build contextual representations** of a **word's meaning** that **integrate information from surrounding words**, helping the model learn how words relate to each other over large spans of text.





# Attention is all you need **Causal self-attention**

- Causal the model has **access to all of the inputs up to and including the one under consideration**
- In general **bidirectional self-attention**, the context can include future words
  - Bidirectional attention was used by BERT model.



# Attention is all you need **Query, Key and Value**

- Consider the three different roles that **each input embedding plays**
- Query
  - As the **current focus** of attention when **being compared to all** of the other **preceding** inputs
- Key
  - In its role as a **preceding input** being compared to the current focus of attention
- Value
  - used to **compute the output for** the current focus of attention.






# Attention is all you need **Query, Key and Value**

$$\text{SelfAttention}(Q, K, V) = \text{SoftMax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

- Calculate the value  $a_3$  the third element in a sequence using causal self attention

I. Generate key, query, value vectors






# Attention is all you need **Query, Key and Value**

- Calculate the value  $a_3$  the third element in a sequence using causal self attention

2. Compare  $X_3$ 's **query** with all the other **keys**



# Attention is all you need **Query, Key and Value**

- Calculate the value  $a_3$  the third element in a sequence using causal self attention


3. Divide the score by  $d_k$

4. Apply softmax to turn it into weights

5. Weight each value

6. Sum the weighted value vectors

Output  $a_3$

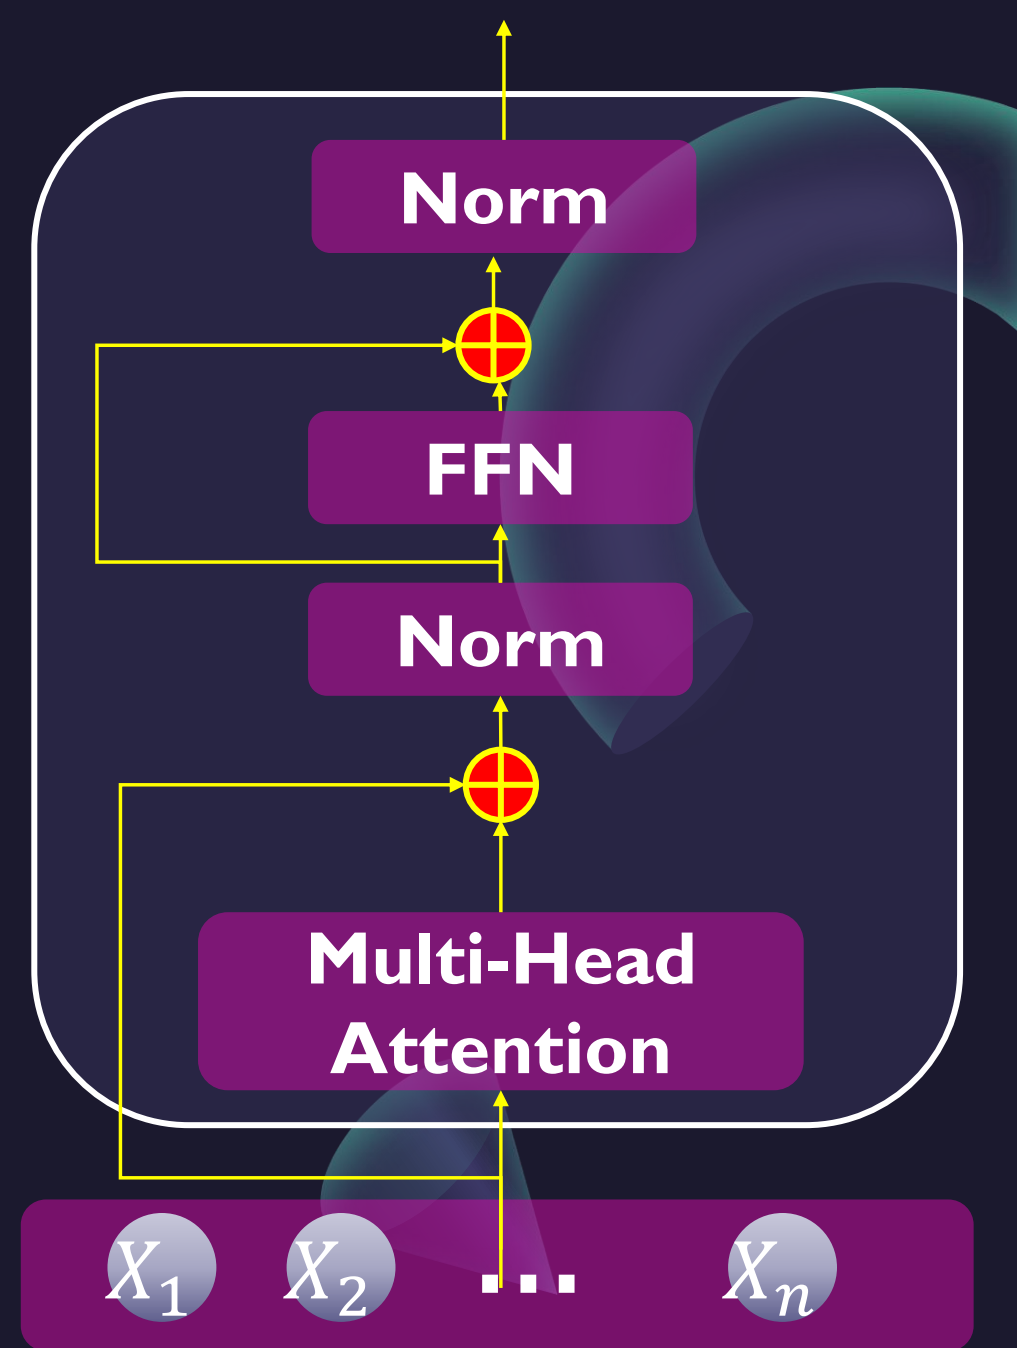


# Transformers **Multi-head Attention**

- Transformers actually compute a more complex kind of attention than the single self-attention
- This is because the different words in a sentence can relate to each other in many different ways simultaneously
- It would be difficult for a single self-attention model to learn to capture all of the different kinds of parallel relations among its inputs
- **multihead self-attention** : sets of self-attention layers, called heads, that reside **in parallel layers at the same depth** in a model, each with its **own set of parameters**. By using these distinct sets of parameters, each head can **learn different aspects of the relationships** among inputs **at the same level of abstraction**.

# Transformers **Blocks**

- includes three other kinds of layers:
  - a feedforward layer
  - residual connections
  - normalizing layers
- $O = \text{LayerNorm}(X \oplus \text{SelfAttention}(X))$
- $H = \text{LayerNorm}(O \oplus \text{FFN}(O))$

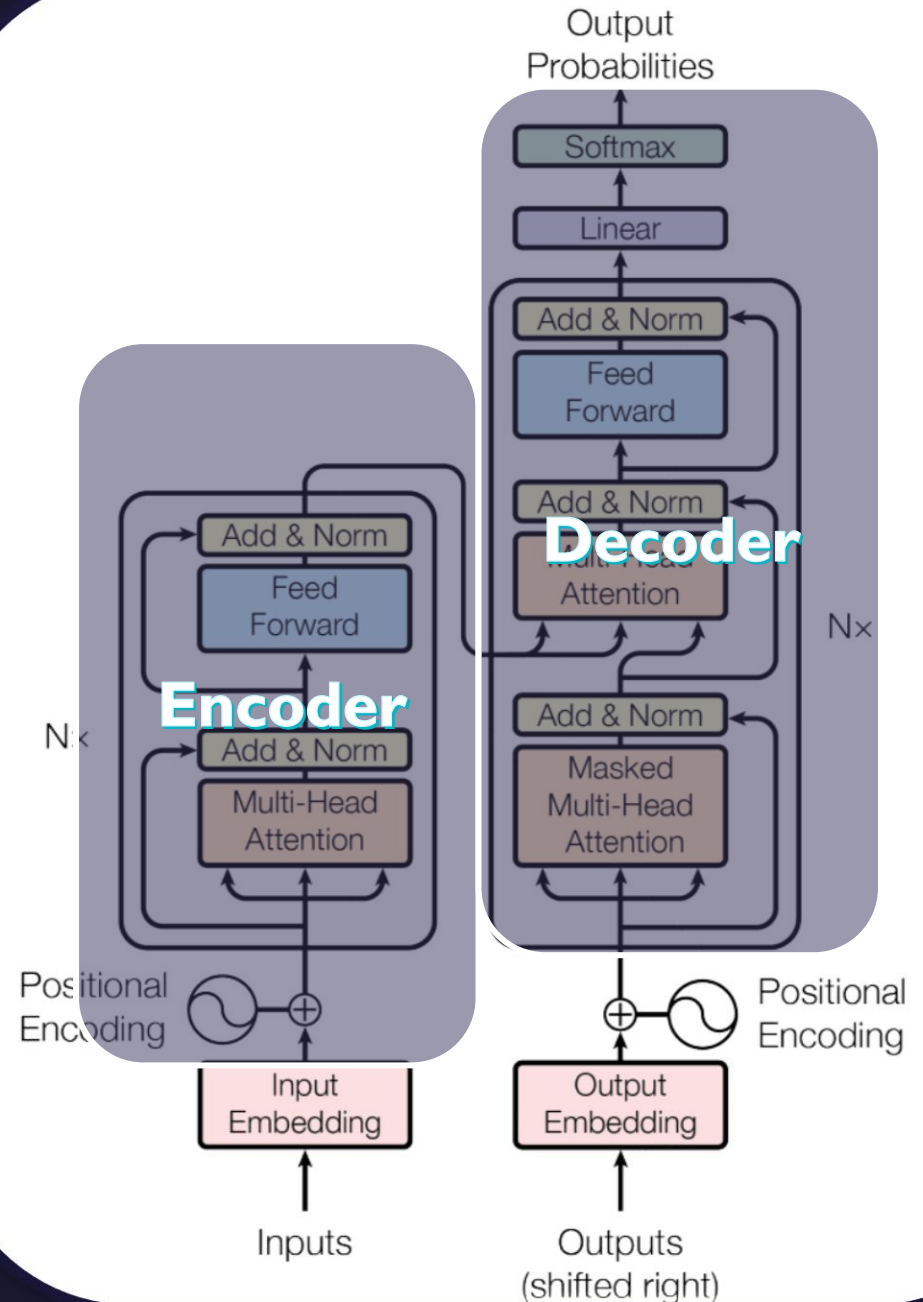




# Transformers

## in “Attention is all you need paper”

- Has two parts “Encoder-Decoder”
- The encoder is composed of a stack of  $N = 6$  identical layers.
- The decoder is composed of a stack of  $N = 6$  identical layers, plus a third sub-layer of multi-head attention that works over the output of the encoder stack



# Transformers

## Encoder/Decoder models

Encoder

BERT

Decoder

GPT

Encoder-  
Decoder



BART



# BERT

## **B**idirectional **E**ncoder **R**epresentations from **T**ransformers


- There are **two** existing strategies for **applying pre-trained** language representations to downstream tasks
- Feature-based
  - Embeddings from Language Models (ELMo)
- Fine-tuning
  - the Generative Pre-trained Transformer (OpenAI GPT)
  - introduces minimal task-specific parameters
  - trained on the downstream tasks by simply fine-tuning all pretrained parameters



# BERT

## **B**idirectional **E**ncoder **R**epresentations from **T**ransformers

- BERT paper claimed that “**The major limitation** is that standard language models are **unidirectional**”
- and this **limits the choice of architectures** that can be used during pre-training
- in OpenAI GPT
  - the authors use a **left-to-right** architecture
  - where **every token can only attend to previous tokens** in the self-attention layers of the Transformer
- Such restrictions are **sub-optimal for sentence-level tasks**





# BERT

## **B**idirectional **E**ncoder **R**epresentations from **T**ransformers

- BERT alleviates the previously mentioned unidirectionality constraint by using
  - “masked language model” (MLM) pre-training objective
  - The masked language model randomly masks some of the tokens from the input
  - and the objective is to predict the original vocabulary id of the masked word based only on its context
- Unlike left-to-right language model pre-training
  - the MLM objective enables the representation to fuse the left and the right context
  - which allows to **pretrain a deep bidirectional Transformer**

# BERT

## Bidirectional Encoder Representations from Transformers

- BERT use also in the pretrain phase “next sentence prediction” task

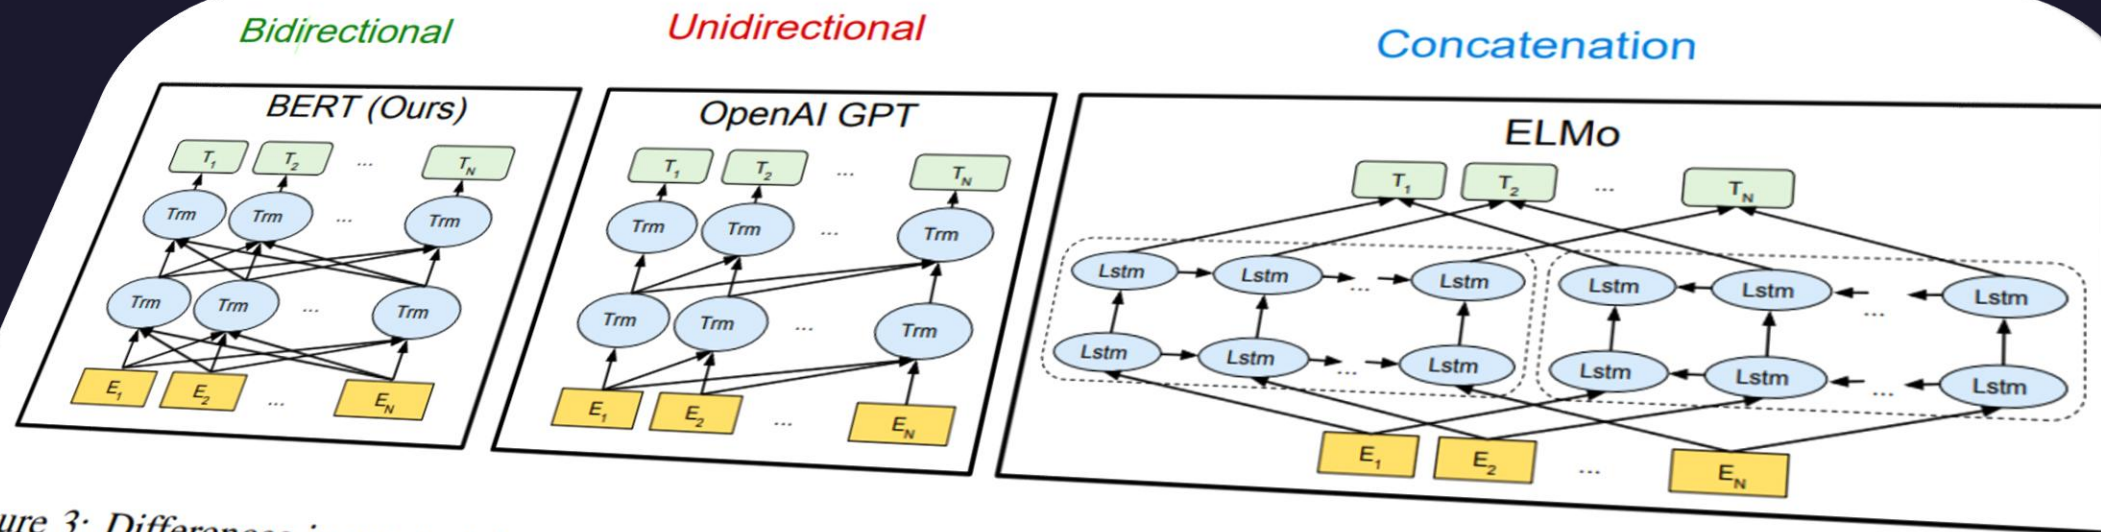





Figure 3: Differences in pre-training model architectures. **BERT** uses a bidirectional Transformer. **OpenAI GPT** uses a left-to-right Transformer. **ELMo** uses the concatenation of independently trained left-to-right and right-to-left LSTMs to generate features for downstream tasks. Among the three, only **BERT** representations are jointly conditioned on both left and right context in all layers. In addition to the architecture differences, BERT and OpenAI GPT are fine-tuning approaches, while ELMo is a feature-based approach.

# BERT

## **B**idirectional **E**ncoder **R**epresentations from **T**ransformers

- A distinctive feature of BERT is its **unified architecture across different tasks**
- BERT Base
  - $L = 12, H = 768, A = 12$
  - **110 M** total parameter
- BERT Large
  - $L = 24, H = 1024, A = 16$
  - **340 M** total parameter



# BERT

## Input representation

- For a given token, its input representation is constructed by summing the corresponding **token**, **segment**, and **position embeddings**

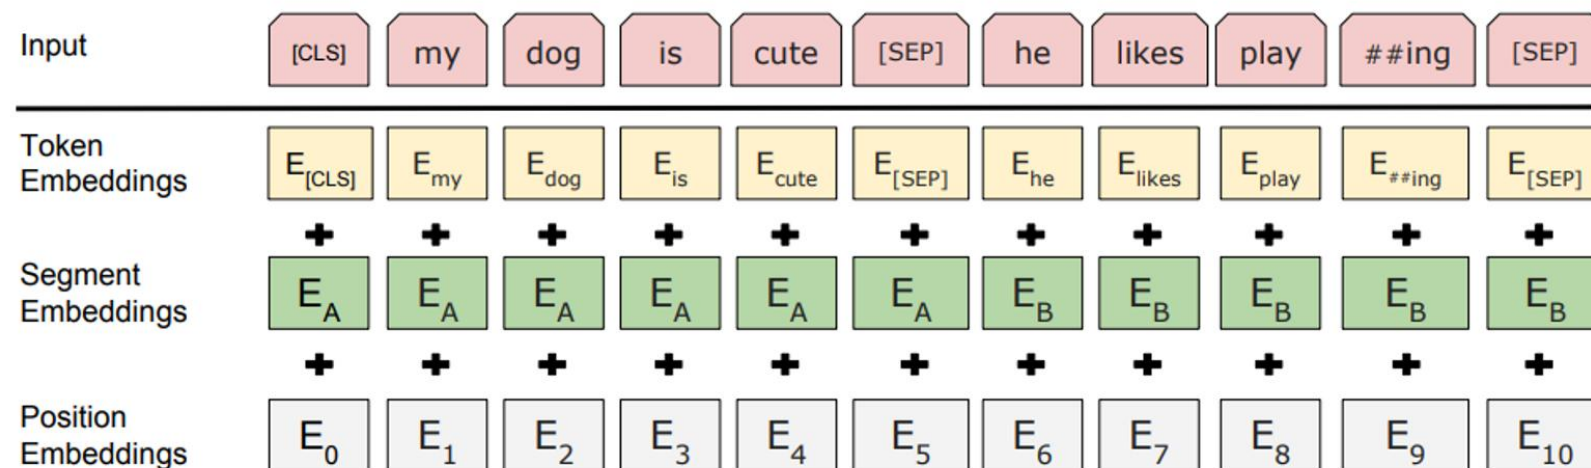


Figure 2: BERT input representation. The input embeddings are the sum of the token embeddings, the segmentation embeddings and the position embeddings.

# BERT

## **Feature-based Approach with BERT**

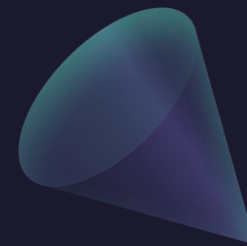

- Not all tasks can be easily represented by a **Transformer encoder architecture**
- There are computational benefits to pre-compute an expensive representation of the training data once and then run many experiments with cheaper models on top of this representation
  - **Fine-tuning all the 110 million or 340 million parameter of a model for every task you want to test the model on is somehow expensive**
  - **In some cases, freezing the model and use it's understanding of language representation on other tasks by simply adding a simple network on top can be so useful and cheaper.**



# BART

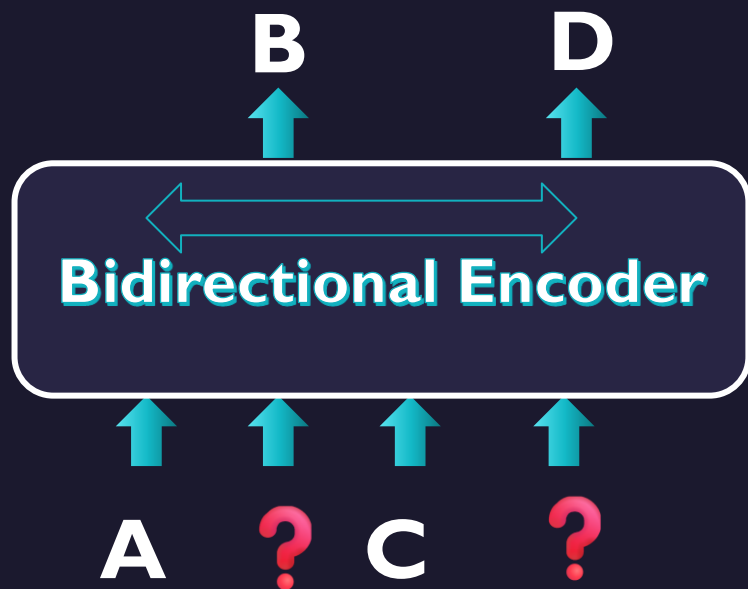
**B**idirectional and **A**uto-**R**egressive **T**ransformers.

- Not all tasks can be easily represented by a **Transformer encoder architecture**
- a denoising autoencoder for pretraining sequence-to-sequence models.
- BART is trained by
  1. corrupting text with an arbitrary **noising** function
  2. learning a model to **reconstruct** the original text
- It use Encoder-Decoder Transformer based architecture
  - can be seen as generalizing to BERT and GPT
  - Bidirectional Encoder & Left-to-Right Decoder



# BART


**B**idirectional and **A**uto-**R**egressive **T**ransformers.



- Random tokens are replaced with masks about 15%
- Missing tokens predicted independently
- BERT **can not easily** be used for generation

# BART


**B**idirectional and **A**uto-**R**egressive **T**ransformers.



- GPT tokens are predicted auto-regressively
- GPT can be used for generation
- But words can only condition on leftward context , and can't learn bidirectional interactions


# BART

**B**idirectional and **A**uto-**R**egressive **T**ransformers.



- BART , inputs to the encoder don't need to be aligned with the decoder , allowing for arbitrary noise functions
- BART is particularly effective when fine tuned for text generation but also works well for comprehension tasks.

# BART Architecture



- BART uses the standard sequence-to-sequence Transformer architecture from “**attention is all you need**”
  - except, following GPT, that they modified **ReLU** activation functions to **GeLUs**
  - initialized parameters from **N (0, 0.02)**.
- The architecture is closely related to that used in BERT with the following differences
  - BERT uses an additional feed-forward network before word prediction , which BART does not
  - Having a decoder that each layer in it preform cross attention over the final hidden encoder
  - BART contains roughly **10%** more parameters than the equivalently sized BERT model.

# BART

## Pre-training

- BART is trained by corrupting documents and then optimizing a reconstruction loss—the cross-entropy between the decoder's output and the original document

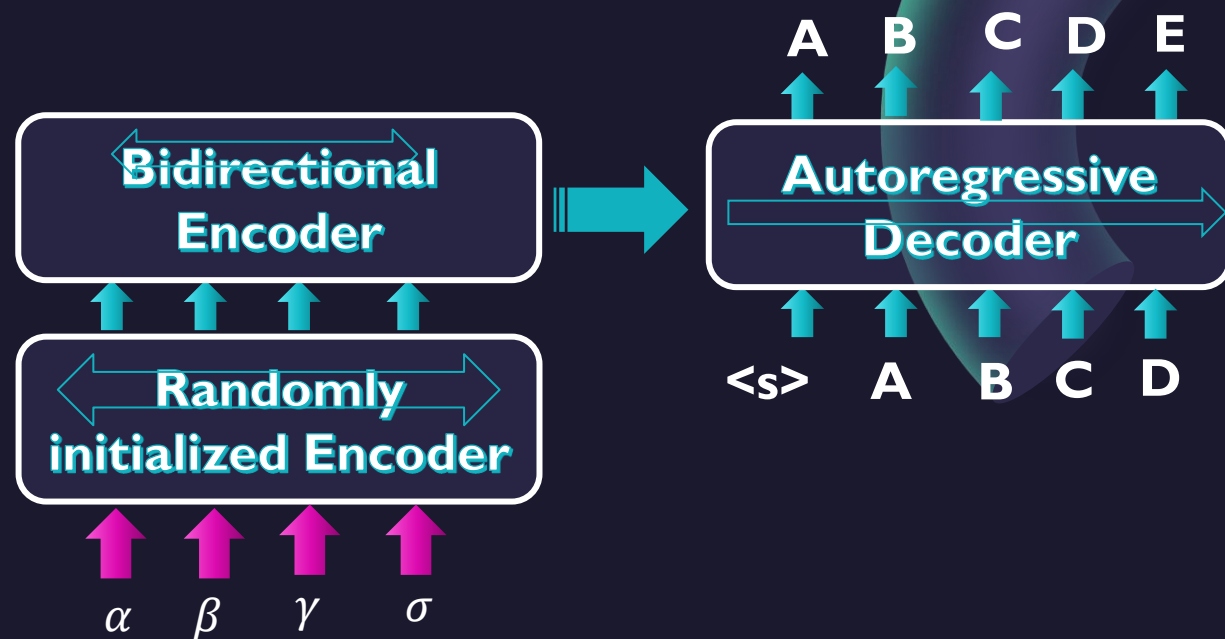




# BART

## In Machine Translation

- It is possible to use the entire BART model (both encoder and decoder) in Machine translation
  - by adding a new set of encoder
- The model is trained end-to-end
- which trains the new encoder to map foreign words into an input that BART can de-noise to English



# See

- <https://poloclub.github.io/transformer-explainer/> [ ✨ visual article]
- <https://bbycroft.net/llm> [ ✨ visual article]
- <https://jalammar.github.io/illustrated-transformer/>
- <https://github.com/jessevig/bertviz/tree/master>
- <https://colab.research.google.com/drive/1hXIQ77A4TYS4y3UthWF-Ci7V7vVUoxmQ?usp=sharing#scrollTo=T3H0qUZvPOP4>
- <https://youtu.be/LPZh9BOjkQs?si=PP4FIKDsoyw0MrOd>
- <https://youtu.be/wjZofjX0v4M?si=RNAHF-USY88pGb4o>
- <https://youtu.be/eMlx5fFNoYc?si=b5HhVEWYZZHofRSj>
- <https://youtu.be/9-Jl0dxWQs8?si=AqMRg793N9GaviCl>
- [https://colab.research.google.com/drive/1JMLa53HDuA-i7ZBmqV7ZnA3c\\_fvtXnx-?usp=sharing#scrollTo=Q3kICzf7LuA9](https://colab.research.google.com/drive/1JMLa53HDuA-i7ZBmqV7ZnA3c_fvtXnx-?usp=sharing#scrollTo=Q3kICzf7LuA9)
- <https://www.kaggle.com/code/alejopaullier/introduction-to-transformers#Introduction-to-Transformers>
- <https://youtu.be/kCc8FmEbInY?si=NSKmXDyaab-hH8xT>
- <https://youtu.be/C9QSpl5nmrY?si=Bu5zrbxzFs6aRpeG>



Thank You